


# Stock Price Prediction

Shailesh Mahto(50540379), Abhiroop(\_\_\_\_), Shivam(\_\_\_\_), Kshitij(\_\_\_\_)

2024-05-10

```
apple <- tq_get('AAPL', from = '2020-01-01', to = '2024-05-01')
apple <- apple|>
  select(date, close) |>
  mutate(day = seq.int(nrow(apple))) |>
  select(day, close) |>
  as_tsibble(
    index = day
  )
autoplot(apple)
```

```
## Plot variable not specified, automatically selected `vars = close`
```



```
# using all data except 1 year for training
train_apple <- apple |> filter(day <= 989)
```

# ETS Model

1. Used Box-cox transformation to stabilize variance
2. Created ETS model and reported autoselected model: ETS(A,N,N)
3. Checked residuals visually and using Ljung-Box test to ensure it is indistinguishable from white noise.
4. Forecasted for the next 100 days and reported accuracy.

```
train_apple %>%  
  features(close, features = guerrero)
```

```
## # A tibble: 1 × 1  
##   lambda_guerrero  
##           <dbl>  
## 1           1.23
```

```
lambda_guerrero = 1.232205  
  
apple_ets <- train_apple |>  
  model(ets = ETS(box_cox(close, lambda_guerrero)))  
  
apple_ets |> report()
```


```
## Series: close  
## Model: ETS(A,N,N)  
## Transformation: box_cox(close, lambda_guerrero)  
## Smoothing parameters:  
##   alpha = 0.9720357  
##  
## Initial states:  
##   l[0]  
## 167.5884  
##  
##   sigma^2: 70.1303  
##  
##      AIC      AICc      BIC  
## 11028.43 11028.45 11043.12
```

```
components(apple_ets) |> autoplot()
```

```
## Warning: Removed 1 row containing missing values or values outside the scale range  
## (`geom_line()`).
```

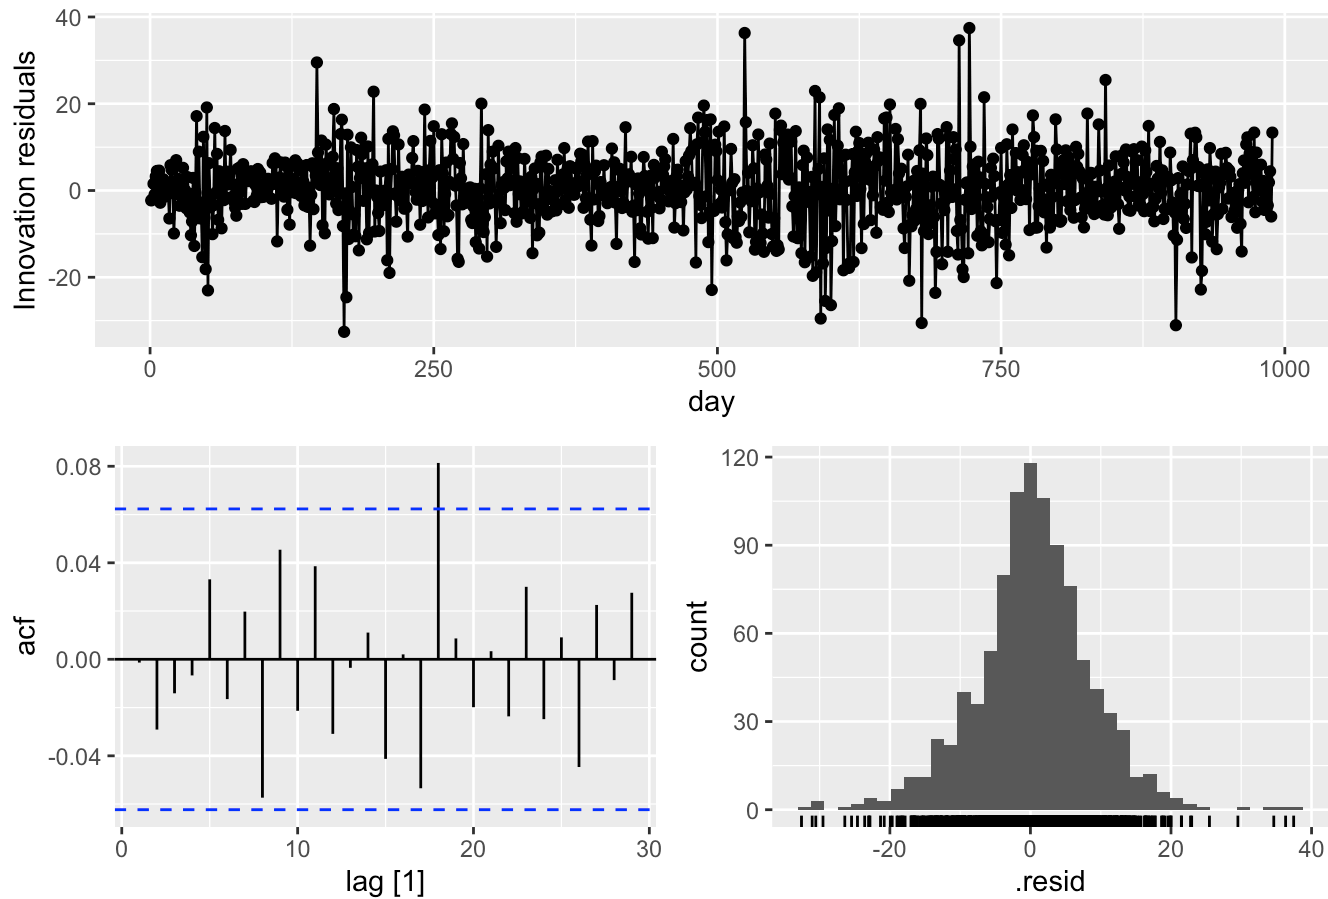
## ETS(A,N,N) decomposition

``box_cox(close, lambda_guerrero)` = lag(level, 1) + remainder`



```
apple_ets |>
  gg_tsresiduals() +
  labs(title = "Residual Diagnostics for ETS(A, N, N)")
```

## Residual Diagnostics for ETS(A, N, N)




```
# using lag=min(2m, T/5) for seasonal data and 10 for non-seasonal data
augment(apple_ets) %>%
  features(.innov, ljung_box, lag = 10)
```

```
## # A tibble: 1 × 3
##   .model lb_stat lb_pvalue
##   <chr>   <dbl>   <dbl>
## 1 ets     8.65     0.565
```

```
apple_ets_fc <- apple_ets %>%
  forecast(h = 100)

apple_ets_fc%>%
  autoplot(apple)
```



```
apple_ets_fc |> accuracy(apple)
```

```
## # A tibble: 1 × 10
##   .model .type    ME  RMSE  MAE  MPE  MAPE  MASE  RMSSE  ACF1
##   <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 ets    Test  -11.8  15.3  12.5 -6.81  7.15  6.27  5.74  0.960
```

## ARIMA Model


1. Used Box Cox transformation to stabilize the variance
2. Converted the data to stationary for modelling
3. Created an ARIMA(2,1,4) model.
4. Ensured its residuals are indistinguishable from white noise-visually and using Ljung-Box test.
5. Reported the model's performance.

```
train_apple %>%
  features(close, features = guerrero)
```

```
## # A tibble: 1 × 1
##   lambda_guerrero
##             <dbl>
## 1             1.23
```

```
lambda_guerrero = 1.232205
```

```
# transform data for constant variance
train_apple %>% autoplot(box_cox(close, lambda_guerrero)) +
  labs(y = "Box-Cox transformed closing price")
```



```
# check for number of seasonal differencing = 0
train_apple %>% features(box_cox(close, lambda_guerrero), unitroot_nsdiffs)
```

```
## # A tibble: 1 × 1
##   nsdiffs
##   <int>
## 1      0
```

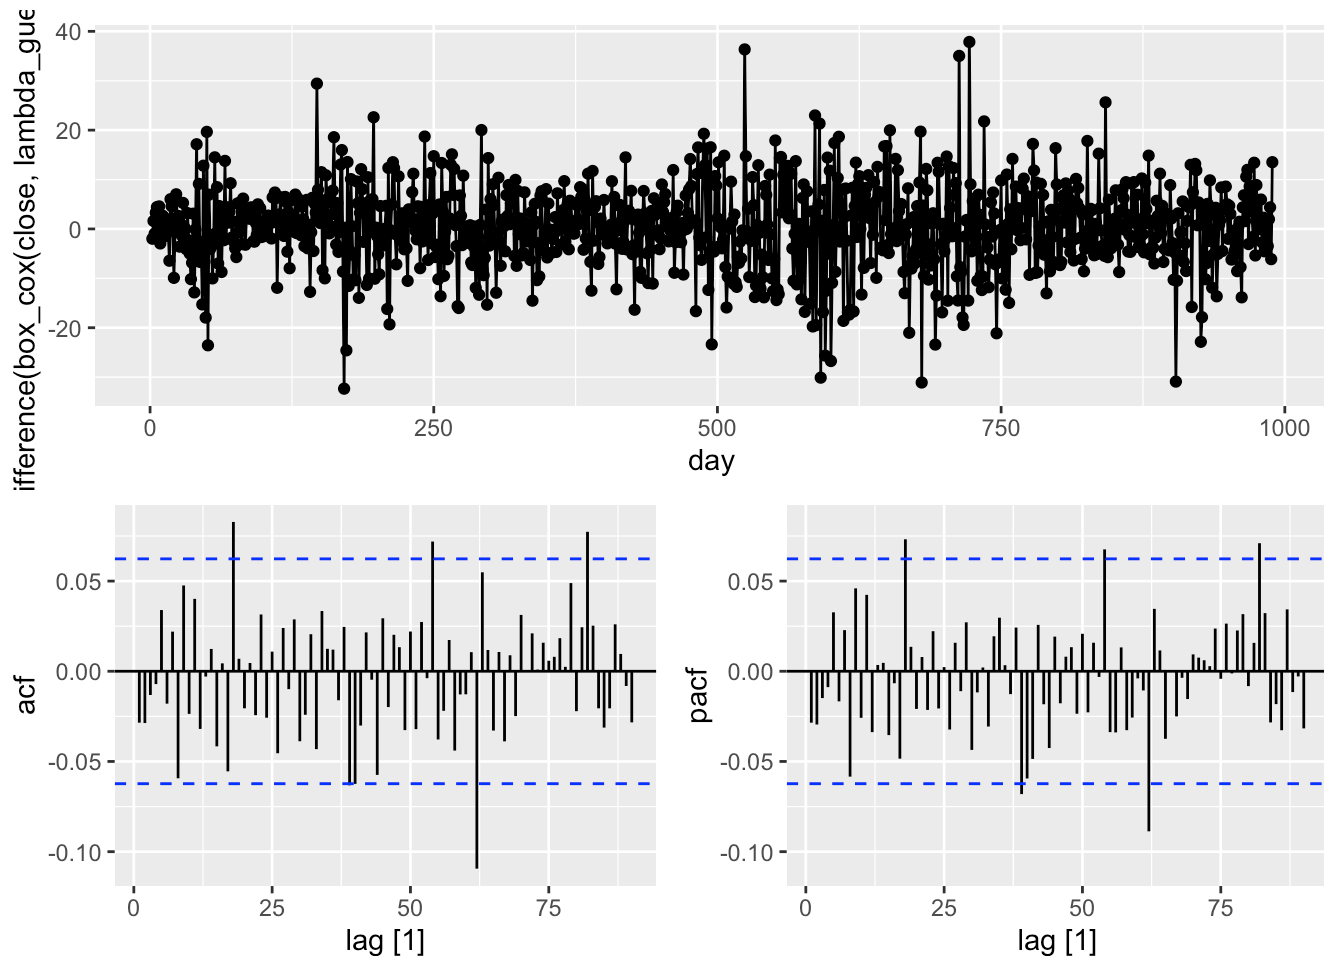
```
# check for first-order differencing = 1
train_apple |> features(box_cox(close, lambda_guerrero), unitroot_ndiffs)
```

```
## # A tibble: 1 × 1
##   ndiffs
##   <int>
## 1      1
```

```
gg_tsdisplay(train_apple, difference(box_cox(close, lambda_guerrero)), plot_type='partial', lag_max = 90)
```

```
## Warning: Removed 1 row containing missing values or values outside the scale range
## (`geom_line()`).
```

```
## Warning: Removed 1 row containing missing values or values outside the scale range
## (`geom_point()`).
```



```
# Obs: ACF drops to 0 quickly, so we can say that the data is now stationary
```

```
train_apple |>
  features(difference(box_cox(close, lambda_guerrero)), unitroot_kpss)
```

```
## # A tibble: 1 × 2
##   kpss_stat kpss_pvalue
##   <dbl>      <dbl>
## 1    0.0483        0.1
```

```
# Obs: The data is stationary
```

```
arima_fit <- train_apple %>%  
  model(  
    arima_auto = ARIMA(box_cox(close, lambda_guerrero), stepwise = FALSE),  
    # arima_auto = ARIMA(box_cox(sale, lambda_guerrero), stepwise = False, approx = False)  
  )  
  
arima_fit |> pivot_longer(everything(), names_to = "Model name",  
  values_to = "Orders")
```

```
## # A mable: 1 x 2  
## # Key:      Model name [1]  
##   `Model name`      Orders  
##   <chr>             <model>  
## 1 arima_auto       <ARIMA(2,1,4)>
```

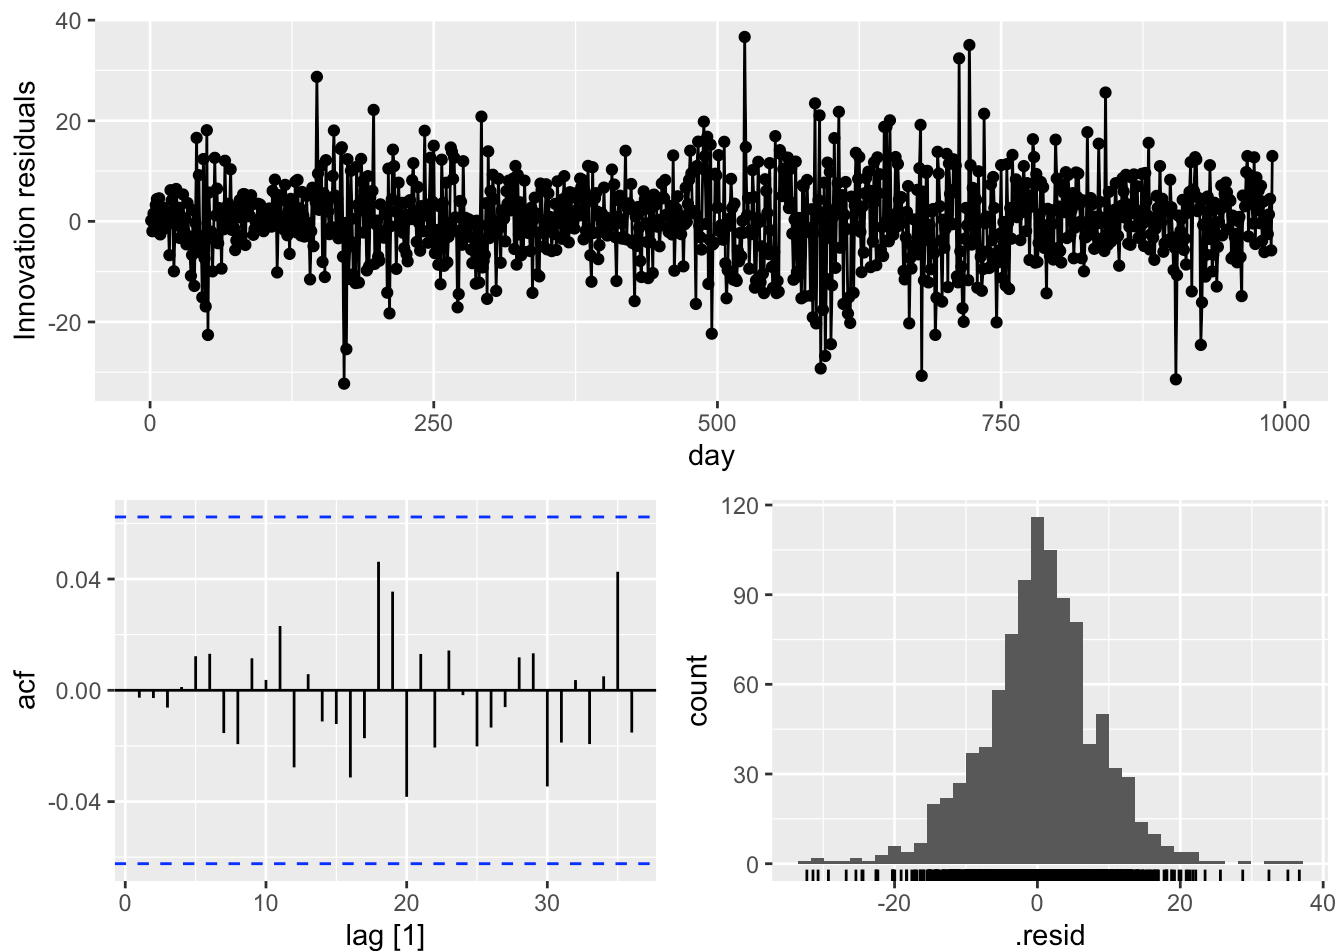
Autoselected model : ARIMA(2,1,4)

```
glance(arima_fit)
```

```
## # A tibble: 1 x 8  
##   .model      sigma2 log_lik   AIC   AICc   BIC ar_roots  ma_roots  
##   <chr>      <dbl>   <dbl> <dbl> <dbl> <dbl> <list>   <list>  
## 1 arima_auto  69.1  -3492. 6997. 6997. 7032. <cpl [2]> <cpl [4]>
```

```
arima_fit |> select(arima_auto) |> gg_tsresiduals(lag=36)
```





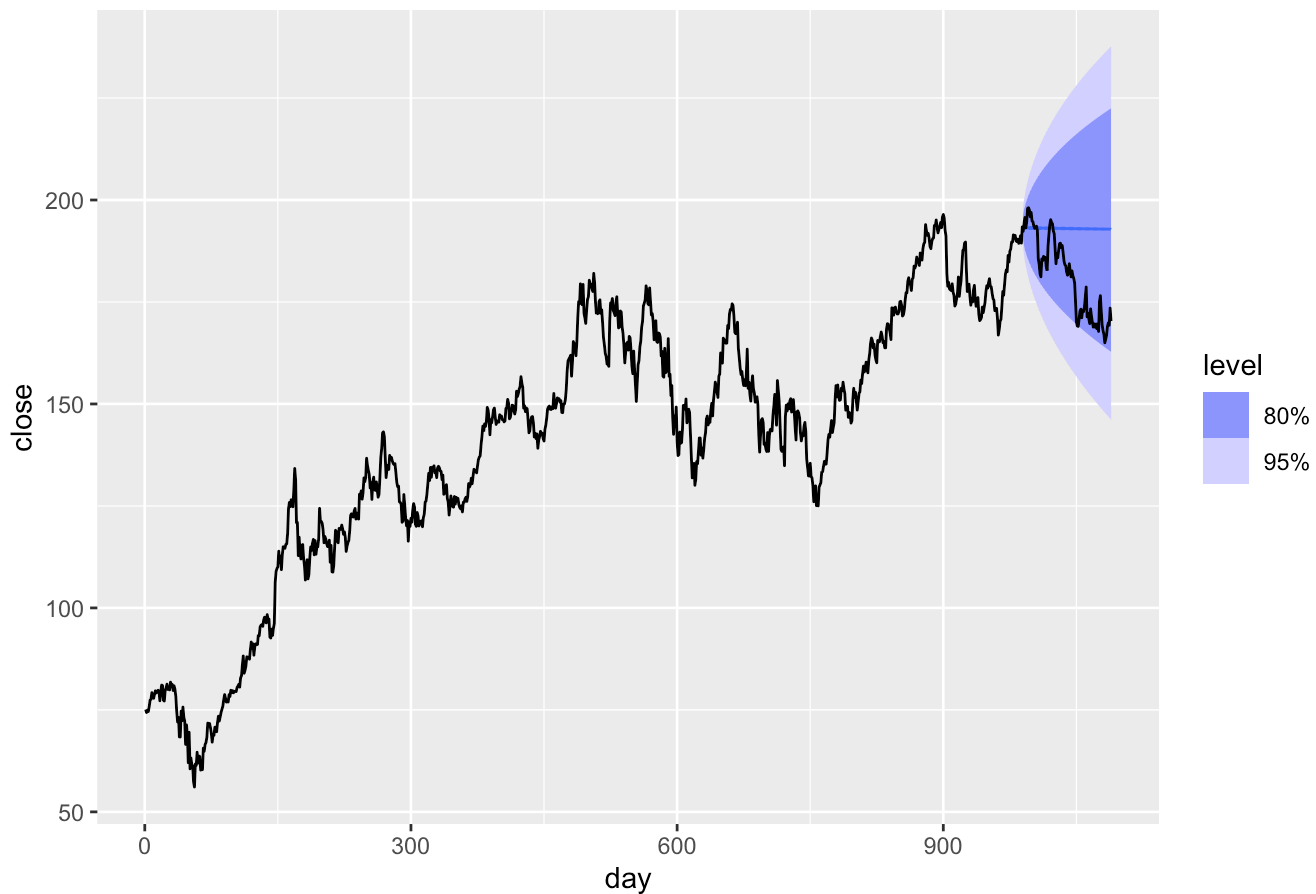
```
augment(arima_fit) |>
  filter(.model == "arima_auto") |>
  features(.innov, ljung_box, lag=24, dof=6)
```

```
## # A tibble: 1 × 3
##   .model    lb_stat lb_pvalue
##   <chr>      <dbl>   <dbl>
## 1 arima_auto  9.76    0.939
```

Observation: The residuals look like white noise, so the model is good. Hence its ready for forecasting.

```
forecast(arima_fit, h=100) |>
  filter(.model=='arima_auto') |>
  autoplot(apple) +
  labs(title = "Closing price actual and forecast using ARIMA(2,1,4)",
        )
```

## Closing price actual and forecast using ARIMA(2,1,4)



```
forecast(arima_fit, h=100) |>
  filter(.model=='arima_auto') |>
  accuracy(apple)
```

```
## # A tibble: 1 × 10
##   .model      .type    ME  RMSE   MAE   MPE  MAPE  MASE  RMSSE  ACF1
##   <chr>      <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 arima_auto Test  -11.7  15.2  12.4 -6.74  7.10  6.22  5.70  0.960
```

## Comparing models

We observe that ARIMA model performs better than the ETS model but their performance is close.

```
train_apple |>
  model(
    arima = ARIMA(box_cox(close, lambda_guerrero) ~ pdq(2,1,4)),
    ets = ETS(box_cox(close, lambda_guerrero) ~ error("A") + trend("N") + season("N"))
  ) |>
  forecast(h = 100) |>
  accuracy(apple)
```

```
## # A tibble: 2 × 10
##   .model .type    ME  RMSE   MAE   MPE  MAPE  MASE  RMSSE  ACF1
##   <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 arima  Test  -11.7  15.2  12.4  -6.74  7.10  6.22  5.70  0.960
## 2 ets    Test  -11.8  15.3  12.5  -6.81  7.15  6.27  5.74  0.960
```